

# Deep Learning - Gesture Recognition

**Submitted by: Sandesh Kandagal and Neelima Boddapati**

## Problem statement (in brief):

As Data scientists, we want to develop a feature in smart-TVs (home electronics) to recognize 5 different gestures performed by the user. This feature will help them to control the TV without using a remote.

Gesture	Control
Thumbs up	Increase the volume
Thumbs down	Decrease the volume
Left swipe	'Jump' backwards 10 seconds
Right swipe	'Jump' forward 10 seconds
Stop	Pause

## Understanding the data:

The training data consists of a few hundred videos categorized into one of the five classes. Each video (typically 2-3 seconds long) is divided into a sequence of 30 frames (images). These videos have been recorded by various people performing one of the five gestures in front of a webcam - similar to what the smart TV will use.

## Data preprocessing:

The data generator will help in preprocessing the images of different dimensions in batches. Using the data generator data preprocessing steps such as image cropping, resizing and normalization were performed.

## Exploring different batch sizes:

Batch sizes: 32, 64, 128 and 256 were explored for model building.

Observations:

- o The kernel repeatedly crashed when using Batch sizes 128 and 256 and hence were not used.
- o All models exhibited better performance with Batch size 32 than 64. Hence, batch size 32 is used in final model evaluation below.

## Modeling:

The following models were built on the data

Conv3D  
Conv2D + RNN LSTM  
Conv2D + GRU LSTM  
Transfer Learning imagenet +LSTM

## Parameters held constant:

Batch size: 32

Epochs: 30

Input shape: 15, 80, 80, 3

- o 15 (frames), 80 (width), 80 (height), 3 (channels)

All models have consecutive layers and with Relu activation function in all layers

Output layer: Dense (5) with Softmax activation function

MaxPooling (2,2)

padding='same'

BatchNormalization

### Model summary:

Model Name	Model	Model parameters	Parameters	Result (Best accuracy metrics)			Observations & Decision
				Epoch	Training	Validation	
Model 1	Conv3D	Hidden layers: 3 layers (16, 32, 64), Kernel size: (3,3,3) FC layer: Dense (512) Dropout: 0.5 Optimizer: <b>Adam</b>	3,350,853	30	93.59	93.75	<b>Observations:</b> The model performs well with good training and validation accuracy. <b>Decision:</b> Explore if changing the optimizer (Model 2) improves accuracy and proceed with model building (Model 3) using the better optimizer.
Model 2	Conv3D	Hidden layers: 3 layers (16, 32, 64), Kernel size: (3,3,3) FC layer: Dense (512) Dropout: 0.25 Optimizer: <b>SGD</b>	3,350,853	24	99.32	81.25	<b>Observations:</b> Changing the optimizer to SGD resulted in overfitting on train data and drop in validation accuracy. <b>Decision:</b> Explore if increasing no of neurons while using Adam as optimizer (Model 3) will improve the accuracy.
Model 3	Conv3D	Hidden layers: 3 layers (32, 64, 64), Kernel size: (3,3,3) FC layer: Dense	3,449,157	24	93.53	81.25	<b>Observations:</b> There is no improvement with increase in no of neurons. <b>Decision:</b> Exploring if

Model Name	Model	Model parameters	Parameters	Result (Best accuracy metrics)			Observations & Decision
				Epoch	Training	Validation	
		(512) Dropout: 0.5 Optimizer: <b>Adam</b>					changing the kernel size (Model 4) will improve model performance.
Model 4	Conv3D	Hidden layers: 3 layers ( <b>16, 32, 64</b> ), Kernel size: ( <b>5,5,5</b> ) FC layer: Dense (512) Dropout: 0.5 Optimizer: <b>Adam</b>	3,606,437	23	74.4	68.750	<b>Observations:</b> There is no improvement with changing the kernel size. <b>Decision:</b> Hence, the best Conv3D model is Model 1. Explore if Conv2d + LSTM will improve model performance.
conv2d 1	CONV2d +RNN LSTM	3 Conv2D groups: (32, 64, 128) Kernel size: (3,3) LSTM: LSTM (128) Dense (64) Dropout: 0.5 Optimizer = Adam	6,917,029	29	81.60	56.25	<b>Observations:</b> This model performance is poor compared to Conv3D models. Validation losses are fluctuating without a concurrent gain in validation accuracy. <b>Decision:</b> Explore if CONV2d +GRU LSTM (Conv2d2) performs better.
conv2d 2	CONV2d +GRU	3 Conv2D groups: (32, 64, 128) Kernel size: (3,3)  GRU (128) Dense (64) Dropout:	5,262,501	29	87.88	75	<b>Observations:</b> The training and validation accuracy is better than the Conv2d1. <b>Decision:</b> Explore if Transfer

Model Name	Model	Model parameters	Parameters	Result (Best accuracy metrics)			Observations & Decision
				Epoch	Training	Validation	
		0.5 Optimizer = Adam					learning will produce better results.
conv2d3	Transfer Learning imagenet +LSTM	Transfer learning using imagenet LSTM (128) Dense (64) Dropout: 0.5 Optimizer = Adam	3,831,877	11	84.48	81.250	<b>Observations:</b> Achieved good accuracy on train and validation data owing to the transfer learning. It has very few trainable parameters (600,965) compared to Model 1 which has similar performance (3,350,629). <b>Decision:</b> Will use this model as the final model for this dataset

### Final decision:

Considering a scenario wherein the company is limited by resources, we suggest to use the Transfer Learning imagenet +LSTM model with model parameters as identified in conv2d3 as the Final model because of it's low trainable parameters and satisfiable performance.

Whereas in a scenario where the company is not limited by resources, we suggest the Convolution 3D model with model parameters as identified in model 1 as the Final model.